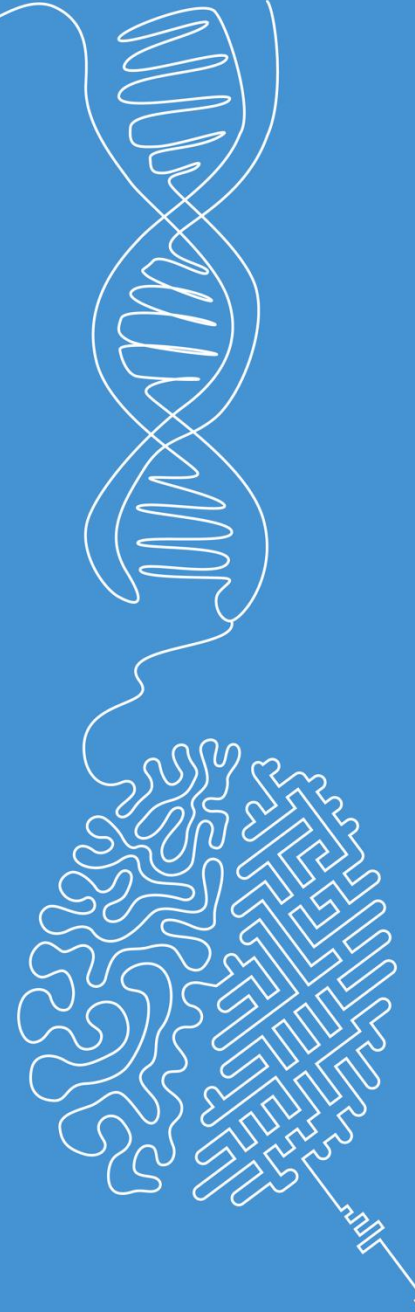


Filtering

Image and Signal Processing

Norman Juchler




Fourier's landscape

We are really wrapping this up now!


Fourier's landscape

Periodic


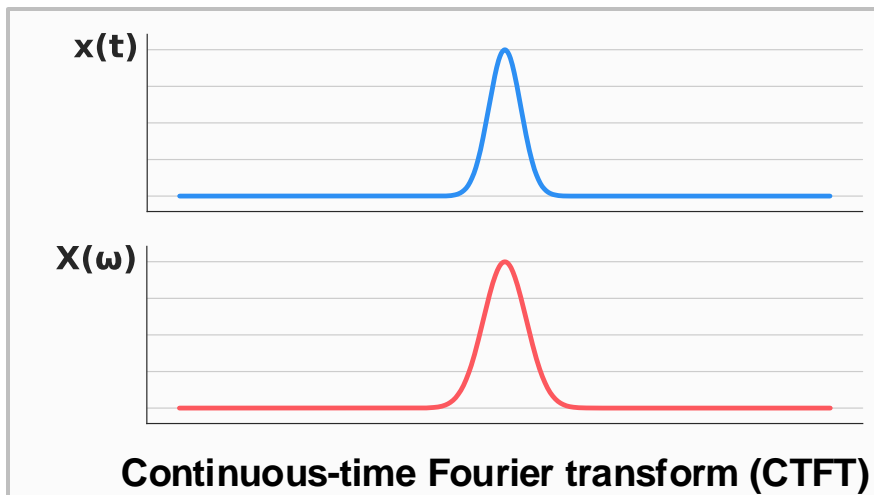
Time-continuous



Time-discrete




Aperiodic



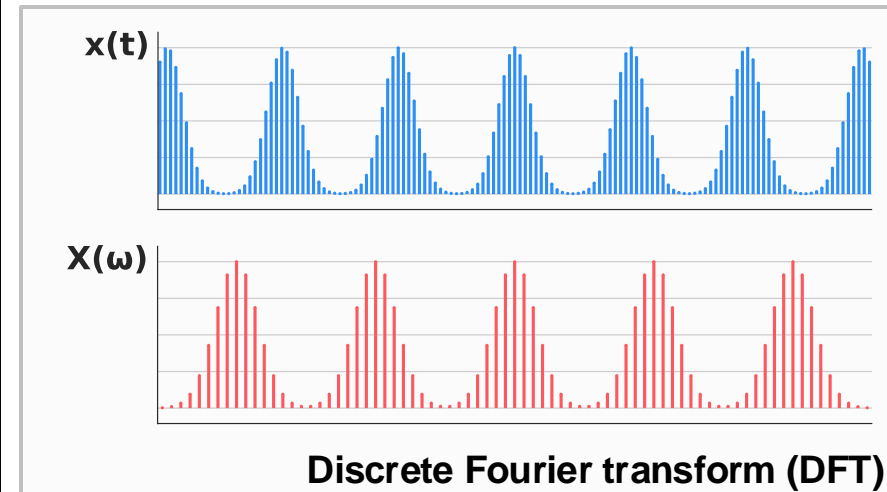
Fourier's landscape

Periodic

Time-continuous





Time-discrete



Related:

- **Fast Fourier transform**
- Discrete cosine transform
- Discrete sine transform

Aperiodic




Generalization:
Laplace Transform

Generalization:
z - Transform

A zoo of different transformations...

- ...but they share many common characteristics:
 - Each has a forward and inverse transformation.
 - They all exhibit duality between time-domain and frequency-domain representations.
 - They share similar mathematical properties, including:
 - Linearity
 - Time/ frequency shifting
 - Convolution theorem
 - ...
- Take-home messages:
 - Fourier theory allows us to decompose signals into their fundamental frequency components.
 - We can seamlessly switch between the time domain and the frequency domain, and vice versa.
 - Different types of Fourier transforms are suited for different types of signals, ensuring the most effective representation.



The discrete Fourier transform

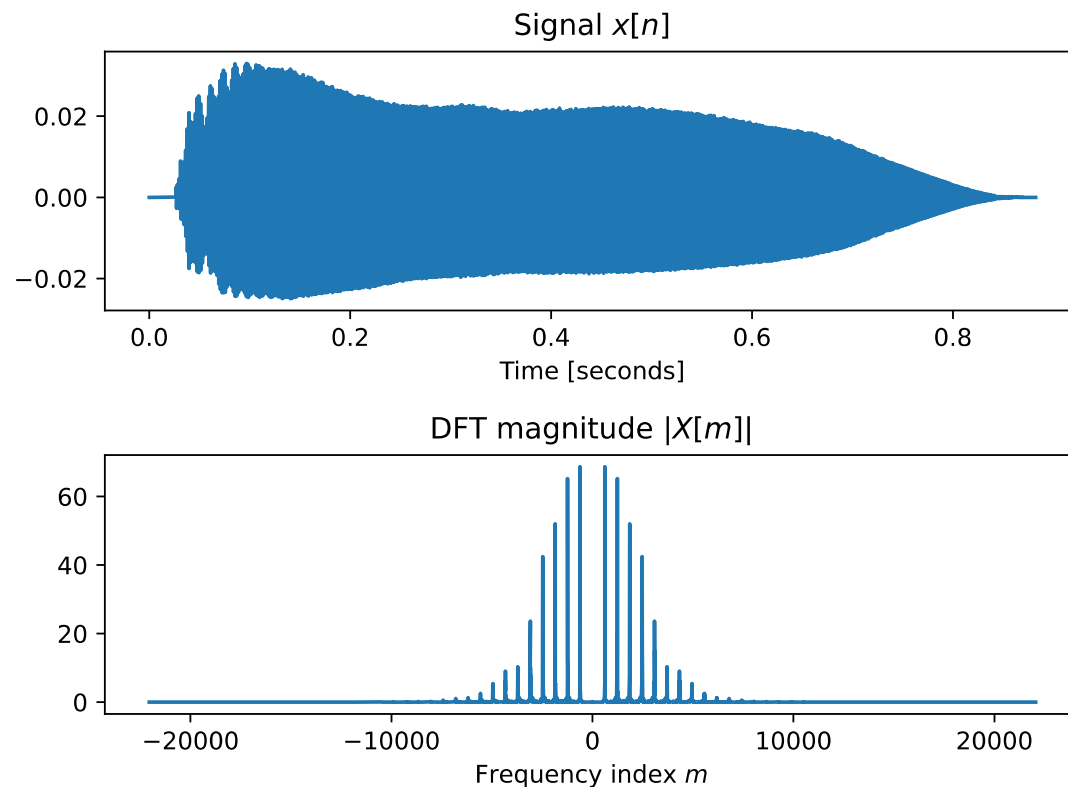
- **Input:** Finite sequence of N equally-spaced samples of a signal
- **Output:** Finite sequence of N equally-spaced samples of the DTFT

$$X[m] = \sum_{n=0}^{N-1} x[n] \cdot e^{-i2\pi \frac{m}{N} n} \quad f_m = \begin{cases} \frac{m}{N} \cdot f_s & \text{if } 0 \leq m < N/2 \\ \frac{m-N}{N} \cdot f_s & \text{if } N/2 \leq m < N \end{cases}$$

- **Observations:**

- The $X[m]$ represents discrete samples of the (continuous) discrete-time Fourier transform
- The $X[m]$ in general are complex numbers (with magnitude and phase)
- The DFT $X[m]$ is periodic with period N (a fact leveraged by `fftshift()`...)
- Every $X[m]$ is associated with the frequency f_m
 - Note that we usually consider frequencies f_m that are negative for $m \geq N/2$ (periodic extension)
 - The first value $X[0]$ is known as the **direct current** (DC) or static component, with $f_0 = 0$
 - $X[0]$ always equals the sum of the sample values

Example: A real sound



$$X[m] = \sum_{n=0}^{N-1} x[n] \cdot e^{-i2\pi \frac{m}{N}n}$$

$$f_m = \begin{cases} \frac{m}{N} \cdot f_s & \text{if } 0 \leq m < N/2 \\ \frac{m-N}{N} \cdot f_s & \text{if } N/2 \leq m < N \end{cases}$$

Look at the illustrations and try to answer the following **questions**:

- What is the sampling frequency?
- How many samples?
- What is the static component about? How is it related to the mean value?
- At which frequency is the peak?

Answers:

- Sampling at 44'100Hz
- Samples $0.88 \cdot 44'100 \approx 38'000$
- Static component: first value $X[0] = \text{sum of } x[n] \approx 0$
- Corresponds to the pitch of the note played (here: $D_5\# = 618.8\text{Hz}$)

The discrete Fourier transform – a naive implementation:


```
1 def dft(x):  
2     '''Compute the Discrete Fourier Transform of an input signal x of N samples.'''  
3  
4     N = len(x)  
5     X = np.zeros(N, dtype=np.complex)  
6  
7     # Compute each X[m]  
8     for m in range(N):  
9  
10        # Compute similarity between x and the m'th basis  
11        for n in range(N):  
12            X[m] = X[m] + x[n] * np.exp(-2j * np.pi * m * n / N)  
13  
14     return X
```

Question: How long does it take to compute the DFT for a signal x ?


Answer: Line 12 needs to be calculated N^2 times, or using the big-O notation: The complexity of this DFT is $O(N^2)$

The Fast Fourier Transform (FFT)

- The FFT is an efficient algorithm to compute the DFT of a discrete signal
- History:
 - Some ideas of FFT can be traced back to C.F. Gauss (1805).
 - The algorithm was first published in 1965 by J. Cooley and J. Tukey.
 - The most popular version is called radix-2 Cooley-Tukey algorithm.
- Key idea: **Divide and conquer**
 - Exploit DFT symmetry to recursively break the computation into smaller sub-problems, eliminating redundant calculations.
 - The roots of unity play a central role here...
 - Instead of directly computing the DFT for $N = N_1 N_2$, it reformulates the problem in terms of for N_1 smaller problems of size N_2 .
 - This approach reduces computational complexity from $O(N^2)$ for a naive implementation to $O(N \log N)$!





The increase of the amount of work (that is, the number of computations) increases with the problem size



A discrete signal of $N=32$ samples (top row) is divided into its even and odd samples (middle-left and middle-right).


Recommended videos

- Reducible / 2023: The Fast Fourier Transform (FFT): Most Ingenious Algorithm Ever? [Link](#)
- Veritasium / 2023: The Most Important Algorithm Of All Time. [Link](#)




The Fast Fourier Transform (FFT)

- The package [scipy.fft](#) offers fast implementations of various FFT algorithms
- See [this tutorial](#) for an overview how to use FFT.
- The main functions:
 - `fft()`, `fft2()`, `fftn()`: Compute the FFT on 1D, 2D, and nD input.
 - `ifft()`, `ifft2()`, `ifftn()`: Compute the inverse DFT in 1D, 2D or nD
 - `rfft()`, `rfft2()`, `rfftn()`: Compute the FFT on real input, faster than the `fft*()` functions.
 - `irfft()`, `irfft2()`, `irfftn()`: Compute the inverse real DFT in 1D, 2D or nD
 - `fftshift()`: `fft*()` and `fftfreq()` return all components at positive frequencies, followed by components at negative frequencies. This function swaps the two halves so that the zero-frequency component is in the center: $[0, 1, 2, -3, -2, -1] \rightarrow [-3, -2, -1, 0, 1, 2, 3]$.
 - `fftfreq()`, `rfftfreq()`: Get the frequencies of each bin returned by the FFT functions, such as the X-axis.




`fftshift()`

→




The short-time Fourier Transform (STFT)

- The DFT makes use of all samples in the signal
- Implicitly, this is assuming that frequency content is “stationary” over the duration of the signal.
- For time-varying signals, this is often not the case
- **Idea:** Divide a long signal up into short pieces (so-called **frames**), analyze each piece separately.
- Parameters:
 - Frame length (number of samples per frame)
 - Hop length (number of samples between two frames)
- STFTs can be visualized with **spectrograms**
 - Constructed by stacking the frames horizontally
 - Hint 1: Use the decibel scale for better visualization
 - Hint 2: Use windowing / apodization to avoid banding artifacts



Spectrogram with linear magnitude scaling (top) and with decibel scaling:
 $A_{\text{dB}} = 20 \cdot \log_{10} A$



Spectrogram without (top) and with windowing (bottom). Use a window (here: Hann) to attenuate discontinuities at the ends of the frames.

Convolution



Revisited

Convolution of continuous-time signals

- The convolution is an operation that combines two functions to produce a third function.

$$y(t) = (h * x)(t) := \int_{-\infty}^{\infty} h(\tau)x(t - \tau) d\tau$$

- In words, integral computes the area under the product of $h(\tau)$ and $x(t - \tau)$ as t varies.
- [Demo applet](#)
- **Video:** 3blue1brown / 2023. But what is a convolution? [Link](#)




Convolution of discrete-time signals

- There exists also a discrete-time definition for convolution:

$$y[n] = (h * g)[n] = \sum_{m=-\infty}^{\infty} h[m]x[n-m]$$

- In the animation right is $x[n]$ the input signal, $h[n]$ the kernel, and $y[n]$ the output signal.



Naive implementation of a discrete-time convolution

```
# Compute the lengths of our input and filter
N = len(x)
K = len(h)

# Allocate an output buffer
y = np.zeros(N)

# Iterate over delay values
for k in range(K):

    # For each delay value, iterate over all sample indices
    for n in range(N):

        if n >= k:
            # No contribution from samples n < k
            y[n] += h[k] * x[n-k]
```

- [scipy.signal.convolve\(\)](#) is a fast alternative
 - Makes use of FFT!
 - Question: Why again?

Convolution and Fourier transform

- **Result:** The Fourier transform simplifies the convolution operation to a simple multiplication in the frequency domain!

$$\begin{aligned} y(t) &= h(t) * x(t) & \Leftrightarrow & Y(\omega) = H(\omega) \cdot X(\omega) \\ y[n] &= h[n] * x[n] & \Leftrightarrow & Y[m] = H[m] \cdot X[m] \end{aligned}$$

- Computing the convolution in the frequency domain is very easy!
- We can efficiently recover the time-domain signal via inverse FFT
- Recommended procedure (for discrete signals):
 - Pad $x[n]$ and $h[n]$ to the same length (if necessary)
 - Use the FFT to compute DTFs $X[m]$ and $H[m]$
 - Multiply $X[m]$ and $H[m]$ elementwise
 - Compute the inverse DFT to recover $y[n]$

Signal filtering

First concepts

Example: Signal smoothing using convolution

Input:

- Noisy signal $x(t)$
- A filter window $h(t)$


Output

- Filtered signal $y(t)$

Solution:

```
import numpy as np
from scipy import signal

noise = 0.05
width = 50
win = signal.windows.boxcar(width)
x = np.repeat([0., 1., 0.], 100)
x += np.random.normal(0, noise, x.shape)
y = signal.convolve(x, win, mode='same') / sum(win)
```



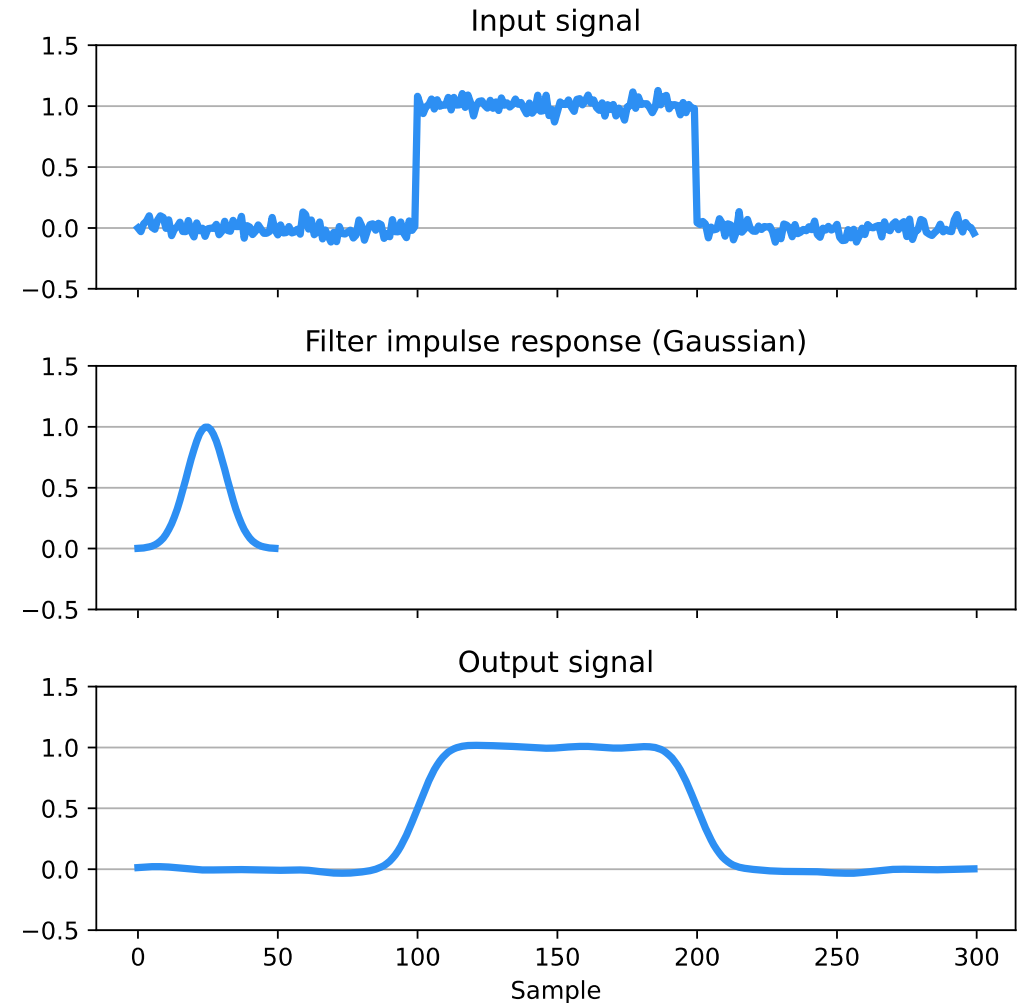
Example: Signal smoothing using convolution

■ Observations:

- We sacrifice some of the signal, but
- We can remove the noise.

```
import numpy as np
from scipy import signal

noise = 0.05
width = 50
win = signal.windows.gaussian(width, std=7)
x = np.repeat([0., 1., 0.], 100)
x += np.random.normal(0, noise, x.shape)
y = signal.convolve(x, win, mode='same') / sum(win)
```




Example: Signal smoothing using convolution

- Many different window types are possible:

- Box
- Triangle
- Blackman
- Hamming
- Hann
- Parzen
- Exponential
- Tukey
- Lanczos
- ...

```
import numpy as np
from scipy import signal

noise = 0.05
width = 50
win = signal.windows.triang(width)
x = np.repeat([0., 1., 0.], 100)
x += np.random.normal(0, noise, x.shape)
y = signal.convolve(x, win, mode='same') / sum(win)
```




Example: Signal smoothing using convolution

- What type of smoothing window should be used? Which parameters to select?
 - A typical filter design question!
 - Requires an understanding of signal characteristics
 - Spectral analysis of signal and filter!









```
import numpy as np
from scipy import signal

noise = 0.05
width = 50
win = signal.windows.lanczos(width)
x = np.repeat([0., 1., 0.], 100)
x += np.random.normal(0, noise, x.shape)
y = signal.convolve(x, win, mode='same') / sum(win)
```



Example: Signal smoothing using convolution

■ Comparison of different windows



$$h_{box}[n] = 1$$

$$h_{gauss}[n] = e^{\frac{1}{2}\left(\frac{n}{\sigma}\right)^2}$$

$$h_{lanc}[n] = \text{sinc}_{\pi} \left(\frac{2n}{M-1} - 1 \right)$$

$$h_{hann}[n] = 0.5 - 0.5 \cos \left(\frac{2\pi n}{M-1} \right)$$


What is the theoretically best low-pass filter?

- Alternative approach: start in the frequency domain
- Goal:** Get rid of all frequencies above some cut-off f_c
- Idea:** Let's define the ideal low-pass filter:




$$H[m] = \begin{cases} 0, & \text{if } f_m > f_c \\ 1, & \text{if } f_m \leq f_c \end{cases}$$

We then can apply the filter in the frequency domain by multiplying it with the signal $X[m]$ and recover the resulting signal $y[n]$ in the time-domain.

- Such a filter is called **ideal low-pass filter**, or **brick-wall filter**.





$H[m]$: The ideal low-pass filter in the frequency-domain




Problems with the ideal low-pass filter

- Has an infinite representation in the time domain
- Cannot be handled properly with DTF, as the filter lengths are limited
- Observations:
 - We already saw that it is difficult to represent sharp edges in the time-domain using frequency components (we require infinitely high frequencies)
 - The same holds true in the frequency domain: The sharper the filter is in the frequency domain, the more time-domain samples are necessary.
 - Note that the filter requires infinitely many samples at $t < 0$! The ideal filter therefore is a non-causal filter.
- See [here](#) for an illustrative demonstration.




Overview: Types of filters


LOW PASS FILTER




HIGH PASS FILTER



BAND STOP FILTER



BAND PASS FILTER



Frequency bands


- A signal is called **band-limited** if its frequency components are limited within frequencies:

$$f_L \leq f \leq f_H$$

- The **bandwidth** of a signal is defined as the difference between the extreme frequencies:

$$B = \Delta f = f_H - f_L$$

- A strictly band-limited signal does not carry energy at frequencies outside the band limits
- In practice, a signal is considered band-limited if its energy outside of a frequency range is low enough to be considered negligible



Amplitude spectrum of a bandlimited signal. Source of illustration: [Wikimedia](#)